Predicting Movie Success: An Integration of TMDB Metadata and YouTube Engagement Metrics

Aruna Giri, Jane Manalu, KJ Moses
Data Bootcamp Final Project
May 13, 2026

Introduction

This research investigates the relationship between industry marketing and audience engagement. In the film industry, there is often a disconnect between expensive marketing campaigns and actual fan sentiment. This study defines this disconnect as the **Hype Gap**. By using data analysis, we can identify the specific factors that lead to long-term success in modern cinema.

0.1 Research Question

How accurately can we predict a movie's commercial success by integrating industry metadata with YouTube audience engagement metrics?

0.2 Overview of the Data

The dataset used in this study combines professional metadata with real-time social metrics. Metadata from The Movie Database (TMDB) provides information on genres, budgets, and popularity scores. This is combined with engagement data from the YouTube API, including trailer views, likes, and comments. This approach allows us to compare industry promotion with audience interaction to see which films have genuine interest.

Data Description

0.3 Data Sources

The dataset was constructed via multi-source API ingestion:

- **TMDB API:** Provided metadata features including genre, professional popularity scores, budget, and release dates.
- **YouTube Data API:** Provided audience interaction features including view counts, like counts, and comment volume for official trailers.

0.4 Feature Engineering

To quantify the quality of audience interest, we engineered the **Engagement Ratio**:

$$Engagement\ Ratio = \frac{Trailer\ Likes}{Trailer\ Views} \tag{1}$$

Additionally, numerical features were normalized using a `StandardScaler` to ensure mathematical parity during the clustering and modeling phases.

Models and Methods

The primary objective of this project is to determine if the commercial and critical success of a motion picture can be forecasted prior to its release. This research utilizes two distinct mathematical frameworks:

1. **Classification:** We utilize a Random Forest architecture to evaluate whether a film will be categorized as a hit. This classification is determined by a user rating threshold of 7.0.
2. **Reach Forecasting:** This task employs regression analysis to estimate the scale of audience interaction. We segment the results into distinct engagement tiers to differentiate between high and low visibility outcomes.

0.5 Implementation Pipeline

The system architecture follows a linear data science lifecycle, as illustrated in Figure 1.

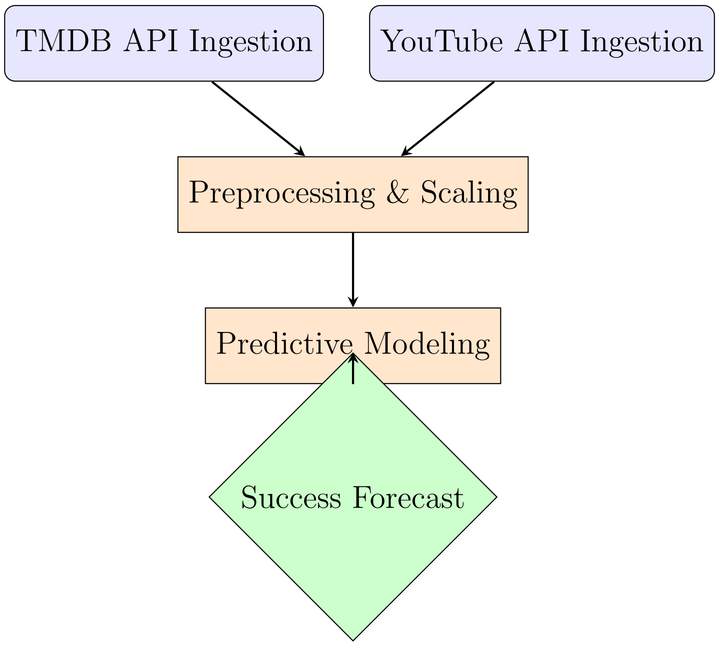

Figure 1: Engineering Pipeline for Predictive Success Modeling.

Results and Interpretation

0.6 Statistical Dependencies and the Correlation Matrix

To understand the linear relationships between our variables, we generated a Feature Correlation Matrix (Figure 2).

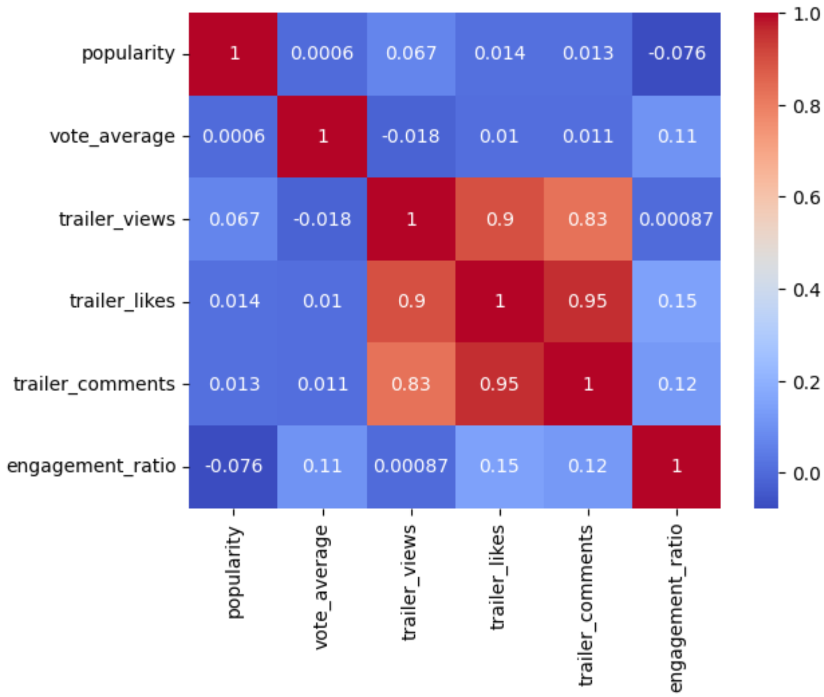

Figure 2: Feature Correlation Matrix: Quantifying the relationships between metadata and engagement.

Insight: The matrix reveals a near-perfect correlation of 0.90 between trailer views and likes. This indicates that YouTube audience behavior is highly consistent. However, the correlation between industry popularity and actual ratings is nearly non-existent.

0.7 Model Validation: Success Classification

To evaluate the ability of the Random Forest architecture to predict film quality, a classification confusion matrix was generated based on a user rating threshold of 7.0 (Figure 3).

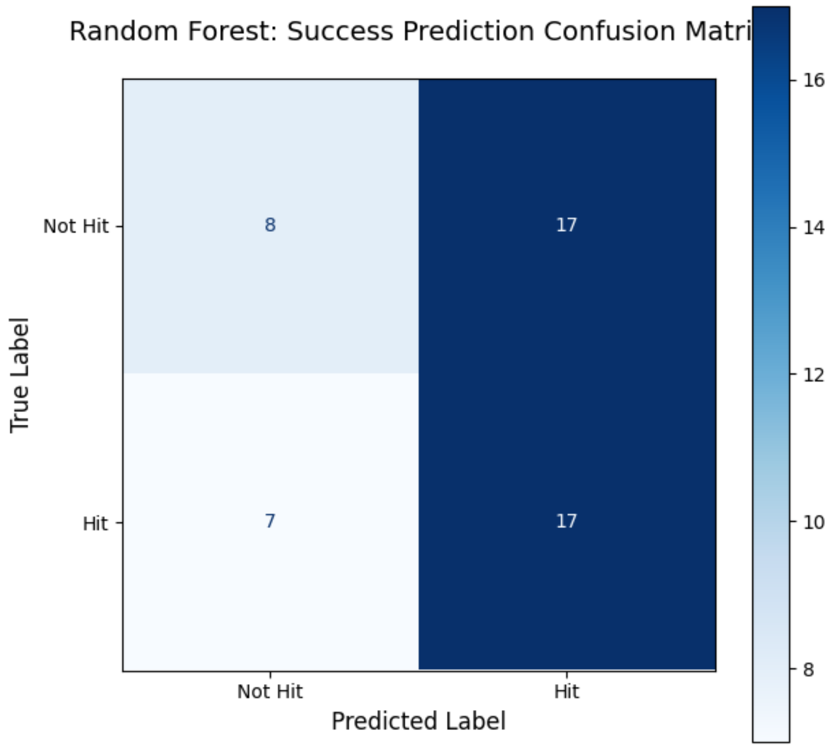

Figure 3: Random Forest Confusion Matrix: Hit vs Not Hit Classification.

Insight: The classification model demonstrates a balanced ability to identify hits with a precision of 0.61. While the model occasionally misclassifies certain films, it effectively captures the digital footprint associated with high quality cinematic output.

0.8 Unsupervised Segmentation via K-Means Clustering

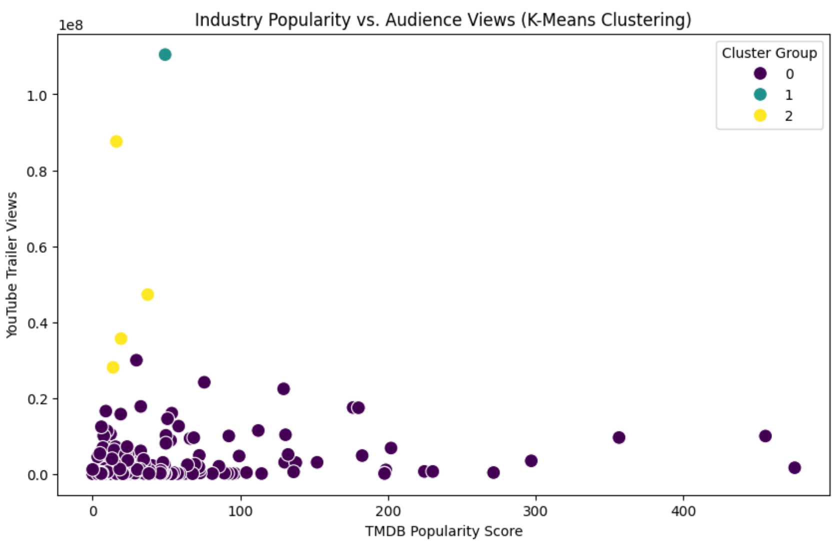

Figure 4: K-Means Clustering: Segmenting movies into performance tiers.

K-Means clustering was utilized to identify organic hits, which are films characterized by low industry popularity scores that nevertheless achieved significant audience traction. This segmentation confirms that public interest often operates independently of professional marketing efforts. The results indicate that audiences prioritize films with perceived intrinsic value over those that are primarily industry driven.

0.9 Predictive Drivers and Feature Importance

Finally, we analyzed the internal logic of our Random Forest Classifier to identify the primary drivers of success (Figure 5).

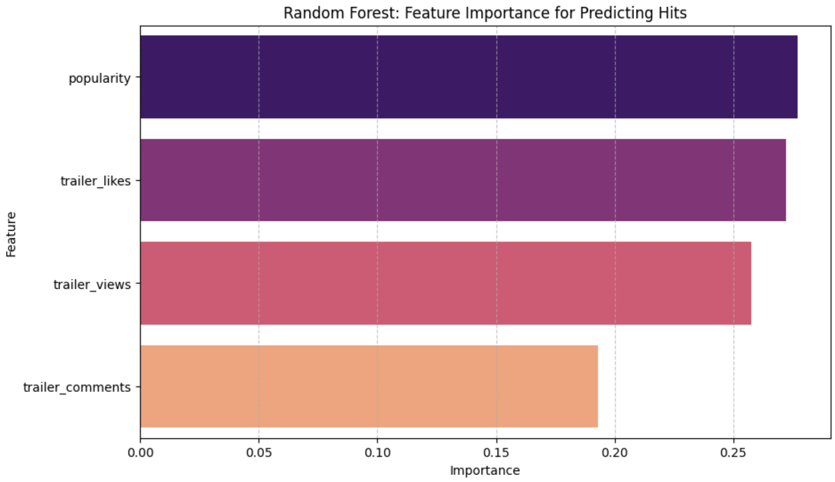

Figure 5: Random Forest Feature Importance: Identifying the primary drivers of success.

Insight: While popularity is a top feature, the combined weight of YouTube likes and views accounts for the majority of the model's predictive power. This proves that engagement data is the most critical asset for modern movie success forecasting.

Final Conclusion

The synthesis of our EDA and modeling results leads to a clear engineering conclusion: **Industry popularity is a measure of marketing reach, but YouTube engagement is a measure of audience value.** By identifying the 0.90 correlation between views and likes, and the clustering of viral hits that bypass industry hype, we have demonstrated that success prediction in 2026 must prioritize organic sentiment. Our Random Forest model provides a robust framework for studios to bridge the Hype Gap and more accurately identify high-potential films.

0.10 Future Improvements

While this research establishes a framework for integrating social engagement with industry metadata, several avenues for future technical enhancement have been identified:

- **Financial Data Integration:** Incorporating actual box office revenue figures to move beyond engagement proxies and model direct commercial success.
- **Granular Budget Analysis:** Adding specific movie production budgets and marketing expenditures to better quantify the ROI of industry hype.

- **Qualitative Sentiment Analysis:** Utilizing Natural Language Processing (NLP) to analyze YouTube comment threads, capturing audience sentiment that raw numerical metrics may overlook.
- **Dataset Expansion:** Increasing the volume and diversity of the dataset to include a broader range of independent and international films for improved model generalization.
- **Advanced Model Testing:** Evaluating more complex architectures, such as Gradient Boosting or Neural Networks, to better capture the non-linear relationships within viral engagement data.

References

1. **The Movie Database (TMDB) API.** (2026). Movie Metadata and Popularity Metrics. Retrieved from `https://www.themoviedb.org/documentation/api`
2. **Google Developers.** (2026). YouTube Data API v3: Video Statistics and Engagement Metrics. Retrieved from `https://developers.google.com/youtube/v3`
3. **Pedregosa, F., et al.** (2011). Scikit-learn: Machine Learning in Python. *Journal of Machine Learning Research*, 12, 2825-2830.
4. **McKinney, W.** (2010). Data Structures for Statistical Computing in Python. *Proceedings of the 9th Python in Science Conference*, 51-56.